

The background of the slide is a dimly lit office interior. A large window on the right side of the frame provides a view of a city skyline at dusk or dawn. The sky is a deep blue with some light clouds. The city features various buildings, including a prominent church with a tall, ornate spire. In the foreground, the office is dark, with some desks and chairs visible. The overall mood is professional and serene.

Booking.com

Water you talking about?

Ivana Rebic & Dennis Bohle | Tel Aviv | 26.02.2019



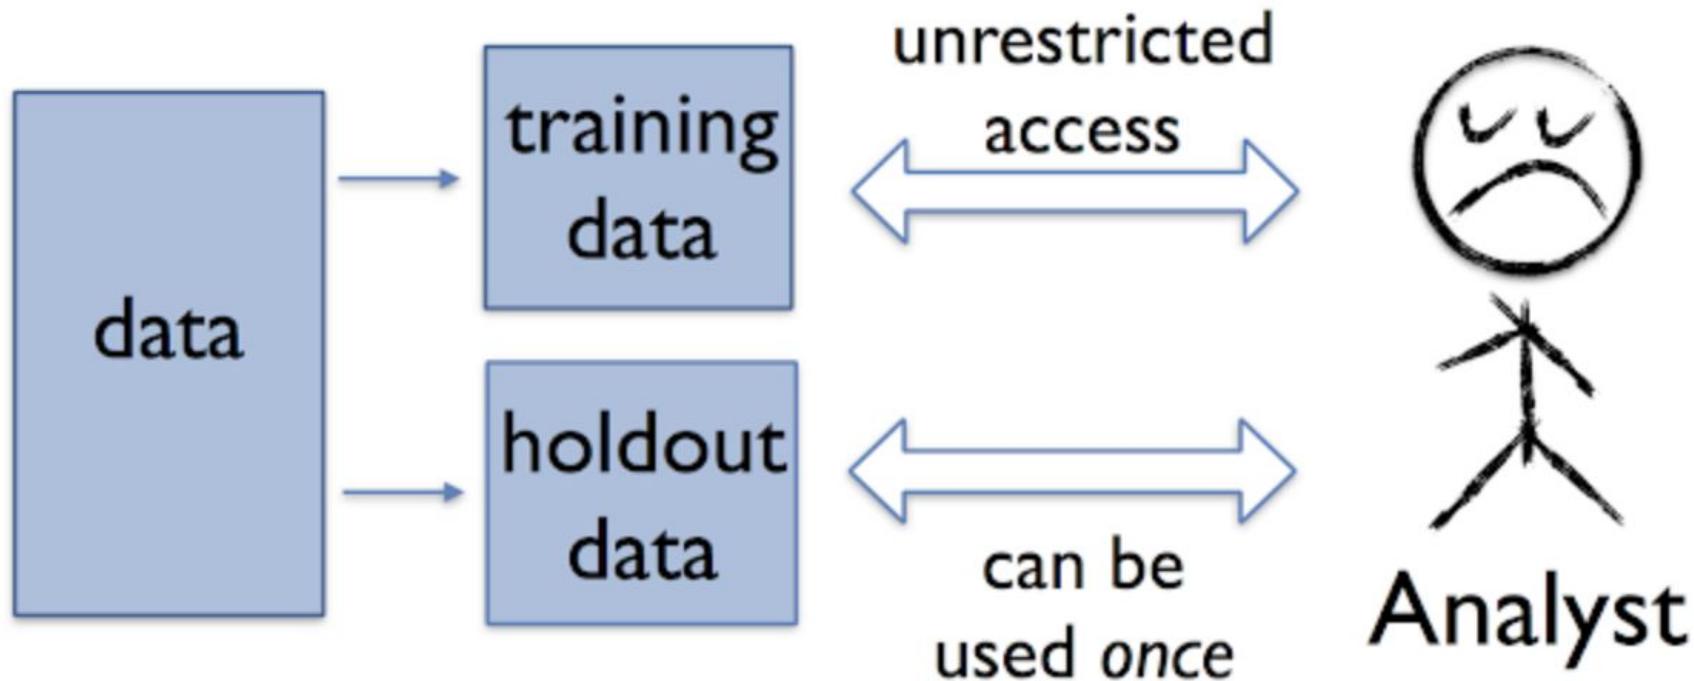
Reusable Holdout Sets

Overfitting on the Kaggle public leaderboard.

Kaggle Iceberg Classification Challenge

| Overview | Data | Kernels | Discussion | Leaderboard | Rules | | |
|----------|---------------|-------------------------|------------|--------------|--------|---------|------|
| # | Δ priv | Team Name | Kernel | Team Members | Score | Entries | Last |
| 1 | — | David & Weimin | | | 0.0801 | 118 | 5d |
| 2 | ▼ 3 | Kohei and Medrr | | | 0.0852 | 114 | 4d |
| 3 | ▼ 105 | Overfitted-Winner 3 | | | 0.0867 | 78 | 4d |
| 4 | — | Mark Rippetoe witnesses | | | 0.0877 | 106 | 5d |
| 5 | ▲ 3 | beluga | | | 0.0880 | 24 | 4d |
| 6 | ▲ 3 | Evgeny Nekrasov | | | 0.0893 | 66 | 4d |
| 7 | ▼ 6 | Pavel Pleskov | | | 0.0921 | 44 | 5d |
| 8 | ▼ 444 | Extreme Overfitted | | | 0.0927 | 13 | 9d |
| 9 | ▲ 3 | AzAkhtyamov | | | 0.0929 | 60 | 4d |
| 10 | ▼ 3165 | Extreme Overfitted | | | 0.0929 | 95 | 5d |

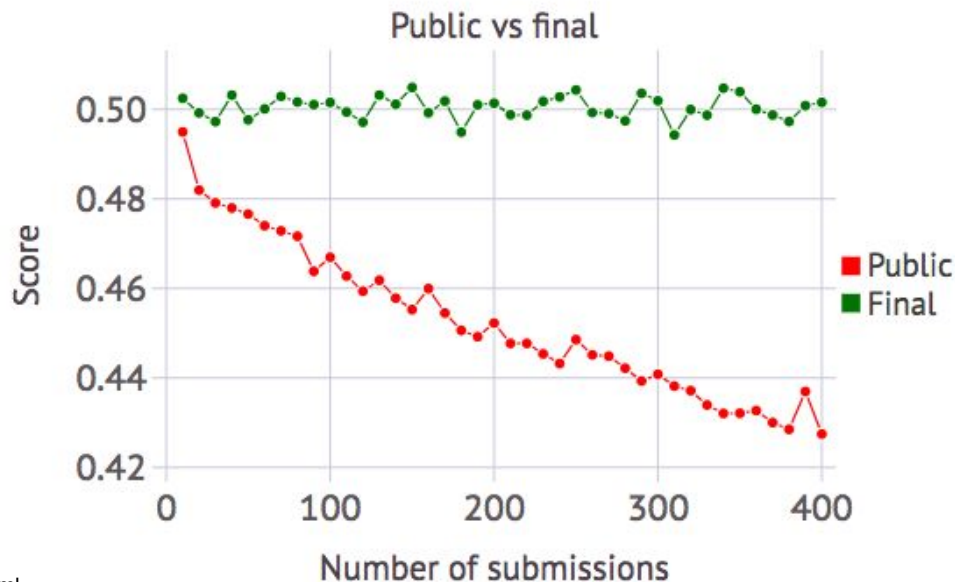
Standard Holdout method.



Wacky Boosting Attack.

Algorithm (Wacky Boosting):

1. Choose $y_1, \dots, y_k \in \{0, 1\}^N$ uniformly at random.
2. Let $I = \{i \in [k] : s_H(y_i) < 0.5\}$.
3. Output $\hat{y} = \text{majority}\{y_i : i \in I\}$, where the majority is component-wise.



The ladder algorithm.

Algorithm 1 Ladder mechanism

Input: Data set S , step size $\eta > 0$

Assign initial estimate $R_0 \leftarrow \infty$.

for round $t = 1, 2, \dots$ **do**

 Receive function $f_t: X \rightarrow Y$

if $R_S(f_t) < R_{t-1} - \eta$ **then**

 Assign $R_t \leftarrow [R_S(f_t)]_\eta$.

else

 Assign $R_t \leftarrow R_{t-1}$.

end if

end for

$$R_S(f) \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n \ell(f(x_i), y_i).$$

The ladder algorithm.


Definition 1 Given an adaptively chosen sequence of classifiers f_1, \dots, f_k we define the **leaderboard error** of estimates R_1, \dots, R_k as

$$\text{lberr}(R_1, \dots, R_k) := \max_{1 \leq t \leq k} |\min_{1 \leq i \leq t} R_D(f_i) - R_t|$$

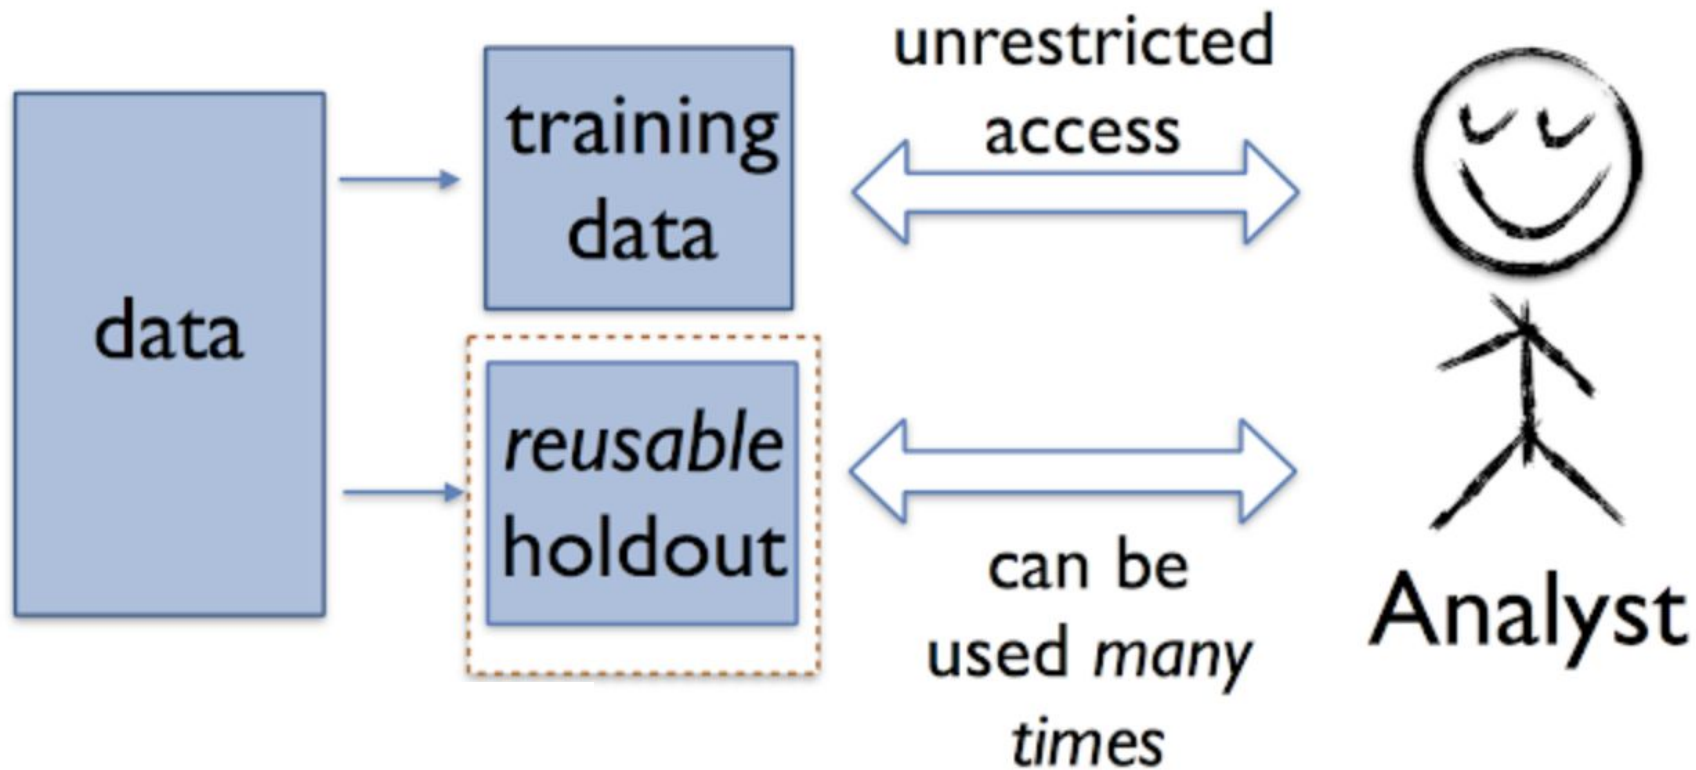
Theorem 1 For any sequence of adaptively chosen classifiers f_1, \dots, f_n there exists η such that with high probability

$$\text{lberr}(R_1, \dots, R_n) \leq \mathcal{O} \left(\frac{\log^{1/3}(kn)}{n^{1/3}} \right).$$

The ladder algorithm under wacky boosting.



Reusable Holdout Method.





Feature Transformation Pipelines.

Model building flow.



Collect.



Collect and clean
the data



Transform.



Apply the
transformations to
the data



Train.



Train a model



Repeat.



Repeat previous
steps until happy



Deploy.



Deploy the model

Model generation at Booking.com.

Most of our data is stored in Hadoop. The most common tools for data extraction and munging are Spark and Hive, but it extends to Python, R and more.

Again for model building we use multiple tools.

Since many of the features are common across models, we have a central place for features that are supported and ready to plug in your model, called **Feature Store**.

Custom feature transformation needed to be implemented separately when model was called in production.



H₂O.ai

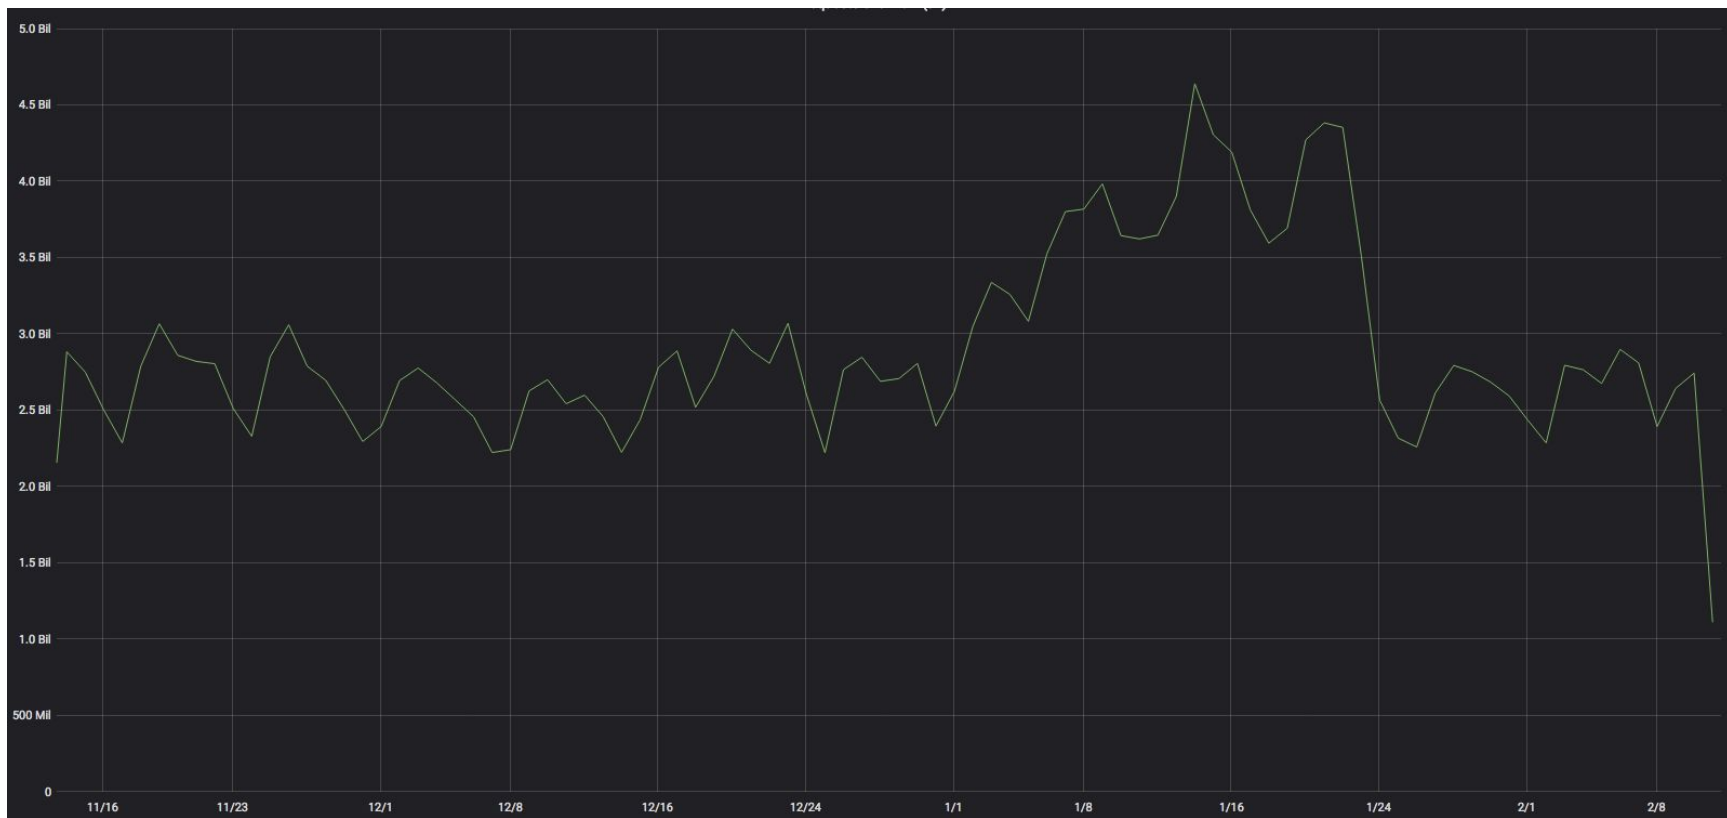
Model serving at Booking.com .

We have model serving platform that provides tooling and monitoring to get model into production.

This is especially useful for models that will be used to predict on live data (for example, as a user searches on the website, or after a booking happens)

It is also a great place to discover existing models, where they are used and what experiments they have been part of.





How to make this faster and more
reliable ?



Introducing Booking Transformations!

We implemented some of the often used transformations as Spark Transformers

Benefits:

- easily serializable
- can work on multiple columns simultaneously
- can be directly deployed
- less offline vs online disparity

Why not Spark MLlib Pipelines ?

Spark comes with the notion of transformers, estimators and pipelines

- ◆ A transformer transforms, . imputation with 0
- ◆ An estimator learns a transformer, imputation with the mean price
- ◆ A pipeline stitches multiple transformations together

What is there not to like about the transformers and estimators that come with spark?

- You can't easily deploy them (but you can store them)
- There is no way to extract the fitted dataframe, so you have to do double computation

Transformations.

Some of the available transformations:

- Imputation
 - {double, string} value imputation
 - {mean, majority} per some other category
- Bucketization
- Interactions
- Concatenations
- Scalers
 - Standardization
 - Normalization
- Target Estimator



Building pipeline.

Each transformation can be created as a separate object, and can then be fed into a transformation pipeline.

A pipeline can then be fit on a dataset to create a pipeline model.

This pipeline model can then be used to transform any dataset.

```
from booking.spark.fv.v1.transformations import Bucketizer, StandardScaler

bucketizer = Bucketizer()\
    .setInputCol("bookingWindow_Look")\
    .setOutputCol("bw_cat")\
    .setBuckets([0.0, 2.0, 5.0])\
    .setInclusionDirection("right")

scaler = StandardScaler()\
    .setInputCol("price")\
    .setOutputCol("price_scaled")

stages = [scaler, bucketizer]

pipeline = Pipeline().setStages(stages)

pipeline_model = pipeline.fit(dataset)

dataset_transformed = pipeline_model.transform(dataset)
```

Frame extractor.

The frame extractor is a simple identity estimator, whose function is to save a pointer to the final, transformed dataset at the end of a pipeline.

You use it by adding it to the end of your list of stages, and then use the `extract()` method to retrieve the transformed dataset

```
from booking.spark.fv.v1.transformations import FrameExtractor

frameExtractor = FrameExtractor()

stages = [scaler, bucketizer, frameExtractor]

pipeline = Pipeline().setStages(stages)

pipeline_model = pipeline.fit(dataset)

dataset_transformed = frameExtractor.extract()
```


Productionizing a pipeline.

The booking transformations are serializable into btl (bookings transformation language), which can be read by our model serving platform and used to create an online version of your transformation pipeline.

```
from booking.spark.fv.v1.transformations import PipelineWriter

path = '.'
writer = PipelineWriter(pipeline_model)
writer.write(path)
```



Benefits.

The booking transformations have a few benefits over the spark provided transformations

1. They do not rely on you creating vectors of your variables
2. They can be performed on multiple columns in parallel
3. They can be productionized using only two lines of code
4. They are ensuring less bugs and feature disparity
5. They shorten the time from model creation to model serving



Coding Target Encoding.

How to use and productionize target estimator.

Using bookings feature transformation pipeline makes it very simple to target encode any desired feature.

All what is needed to do is to specify and add Target Estimator object to the frame extractor!

```
targetEstimator = TargetEstimator() \
    .setInputCol("var1") \
    .setOutputCol("te_var1") \
    .setFoldColumn("fold") \
    .setSlope(10) \
    .setInflectionPoint(50)
```

The inflection point determines how big a sample we still trust.

The slope determines how quickly this trust degrades, by defining the slope around the inflection point.

Note that these are hyper parameters of the encoding that should be tuned when using.

Fold column should already be available in the dataset.

We will create a fold, and apply the transformation on our training dataset.

```
train = train.withColumn("fold", F.round(F.rand()*5))
targetEncoder = targetEstimator.fit(train)

train_te = targetEncoder.transform(train)
```

In order to also apply the target encoder to the test dataset, we have to tell the encoder not to use the folds (there is no danger of leakage in the test dataset)

```
paramList = []
paramList[targetEncoder.holdout] = False

test_te = targetEncoder.transform(train, paramList)
```

If we want to use target estimator, now we just need to add it to the list of our stages with other feature transformations and follow the same steps as mentioned before:

```
from booking.spark.fv.v1.transformations import FrameExtractor


frameExtractor = FrameExtractor()

stages = [scaler, bucketizer, targetEstimator, frameExtractor]

pipeline = Pipeline().setStages(stages)

pipeline_model = pipeline.fit(dataset)

dataset_transformed = frameExtractor.extract()
```

Feature importance and explainability in tree models

Interpreting a complex model: Global vs Local.

Global interpretability

Relationship between the inputs and the dependent variable(s).

Local interpretability

Local interpretations promote understanding of small regions of the conditional distribution.

Game theory.

1. The sum of what everyone receives should equal the total reward
2. If two people contributed the same value, then they should receive the same amount from the reward
3. Someone who contributed no value should receive nothing
4. If the group plays two games, then an individual's reward from both games should equal their reward from their first game plus their reward from the second game



Additive feature attribution method.

Definition 1 *Additive feature attribution methods have an explanation model that is a linear function of binary variables:*

$$g(z') = \phi_0 + \sum_{i=1}^M \phi_i z'_i, \quad (1)$$

where $z' \in \{0, 1\}^M$, M is the number of simplified input features, and $\phi_i \in \mathbb{R}$.

Local accuracy.

Local accuracy states that the sum of the feature attributions is equal to the output of the function we are seeking to explain.

Property 1 (Local accuracy)

$$f(x) = g(x') = \phi_0 + \sum_{i=1}^M \phi_i x'_i \quad (5)$$

The explanation model $g(x')$ matches the original model $f(x)$ when $x = h_x(x')$, where $\phi_0 = f(h_x(\mathbf{0}))$ represents the model output with all simplified inputs toggled off (i.e. missing).

Missingness.

Missingness states that features that are already missing are attributed no importance.

Property 2 (Missingness)

$$x'_i = 0 \implies \phi_i = 0$$

Missingness constrains features where $x'_i = 0$ to have no attributed impact.

Consistency.

Consistency states that changing a model so a feature has a larger impact on the model will never decrease the attribution assigned to that feature.

Property 3 (Consistency) *Let $f_x(z') = f(h_x(z'))$ and $z' \setminus i$ denote setting $z'_i = 0$. For any two models f and f' , if*

$$f'_x(z') - f'_x(z' \setminus i) \geq f_x(z') - f_x(z' \setminus i) \quad (7)$$

for all inputs $z' \in \{0, 1\}^M$, then $\phi_i(f', x) \geq \phi_i(f, x)$.

Shapely Values.

$$\phi_i = \sum_{S \subseteq F \setminus \{i\}} \frac{|S|!(|F| - |S| - 1)!}{|F|!} [f_{S \cup \{i\}}(x_{S \cup \{i\}}) - f_S(x_S)] .$$



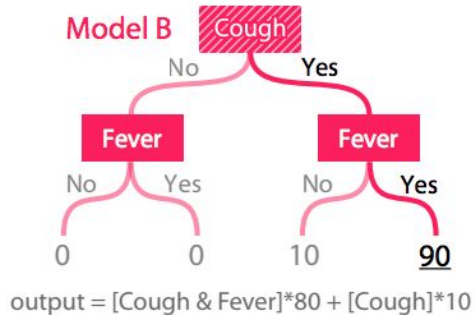
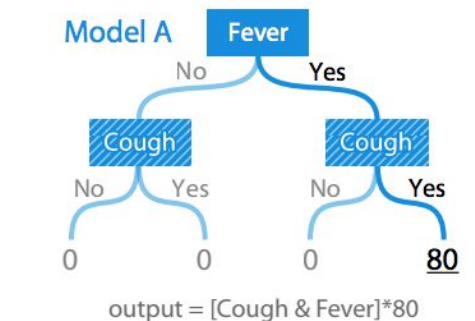
Uniqueness.

Theorem 1 *Only one possible explanation model g follows Definition 1 and satisfies Properties 1, 2, and 3:*

$$\phi_i(f, x) = \sum_{z' \subseteq x'} \frac{|z'|!(M - |z'| - 1)!}{M!} [f_x(z') - f_x(z' \setminus i)] \quad (8)$$


where $|z'|$ is the number of non-zero entries in z' , and $z' \subseteq x'$ represents all z' vectors where the non-zero entries are a subset of the non-zero entries in x' .

Inconsistencies in Gain and Split Count



Model A Attributions

Model B Attributions



Some open research
questions at Booking.com



Just a few examples.

- Session based Information Retrieval
- Interference between A/B-Tests
- Testing under capacity constraints
- Debiasing feedback loops

A modern office interior with a high ceiling, exposed pipes, and large circular pendant lights. In the foreground, a long yellow table sits on a grey stone base. To the left, a long concrete table is surrounded by white chairs, with a woman sitting at it. In the background, several people are working at tables. To the right, there are red and blue modular seating areas. The overall atmosphere is bright and professional.

Thank you!

Booking.com

Questions?

